

COMP 161 - Lecture Notes - 01 - Introduction

January 3, 2016

This course is not your first look at programming. In these introductory notes we look at how this course differs from your experience in COMP160 and how it builds off of your experiences programming in BSL Racket with DrRacket. We also get a brief overview of tools we'll use to craft our programs in this course.

What's a Program?

The title of this course is *Introduction to Programming*. You're probably thinking, "we just had a semester's worth of programming, so what's this?" Well, yes, you know enough about programming to write code to solve pretty much any problem you encounter. However, don't know enough to stay out of trouble. Here's another perspective. In COMP160 you learned about *computation*, most of the fundamental kinds of data that one encounters in computation, and how to bring these things together in a functional manner. The programming you learned focused mainly on describing what you observed or what you hypothesized is going on within a particular computational process. Now, you probably weren't thinking in this way. This is a highly mathematical and scientific perspective.¹ You were probably setting out to write some code, to make a game, to get a programming assignment done. The wonderful thing is that languages like Racket allow you to do both at the same time. They're designed to allow you to focus on the computational process itself. Your programs were as much a *declaration of a computation* as anything else. This is a valid way of attacking the act of programming and you could build an entire career off this way of thinking. So, yes you know how to design and write programs that describe and declare a computational process in a way that allows for Racket to interpret and execute them. But there are other ways of expressing a program.

The word **program** itself implies a plan or a schedule. To see the plan in a Racket program you had to think like the Racket Virtual Machine² and decide the order in which actual computation was enacted. What if instead, you the programmer set the plan? From this we arrive at the perspective³ of programming in which we write procedures *for a computer to carry out and explicitly specifying the order in which they are carried out*. We're now starting to think more about the machine and perhaps less about abstract processes. Programs in this world are called *imperative*. You issue a series of instructions or commands to be carried out over time⁴. If we look deeper into the Racket programming language we could find all the tools we need

¹ hence we call that course *Introduction to Computer Science*

² or "DrRacket"

³ or paradigm

⁴ like a recipe

to program in this fashion with Racket. We will, instead switch gears to a language that more readily supports our new imperative style, C++⁵.

⁵ <http://www.cplusplus.com>

Thus, one of the major goals for this class is to understand this new programming paradigm and the impact of a device centric view of programming. Don't worry, this is by no means starting over. We can and will leverage everything we learned about computation and programming in COMP160. We just need to revisit the fundamental ideas and programming methodologies we learned about programs and computation to fit within our new paradigm and our new language. Before we can do that, we need to find a new set of tools for replace DrRacket.

Programming: Science, Engineering, and Craftsmanship

Programming is part science, part engineering, and part art. You're beginning to understand the science, but its time we think about the art and engineering and take a broader view of the craft of programming. Plumbers, carpenters, and mechanics all have a standard set of tools they use and shared set of goals for the things they do. They all know what good craftsmanship⁶ looks like and how to evaluate the quality of their work and the work of others. The same is generally true of programming. In order to program you need a specific set of tools. When you're programming, you have a specific set of goals you're trying to achieve in order to produce a high quality program. It are these tools and these goals that you need to start thinking about and integrating into your program design process. We should, of course, begin with the goals as they are the end to which programming tools are a means.

⁶ craftsperonship

The Problem

In general, we can say that a program begins well before the first line of code is written. It begins when someone, *identifies a problem* and decides to *solve the problem using computing technology*⁷. Now, sometimes the problems are as straight forward as, "design a game to make me lots of money," and sometime they're less clear like, "How do we enhance the quality of life of people suffering from dementia?" Once someone decides they're going to use computation to model and solve a problem, they must identify the *platforms* on which they'll launch their solution.

⁷ we're just thinking software, but hardware and hybrid systems are options as well

The Platform

When we say platform we generally are thinking about the hardware and operating system layers on which our software runs. However, these days software itself is a viable platform. Some possible platforms are:

- Smartphones: Apple, Android, Windows, Blackberry, etc.
- Wearable, embedded systems: Google Glass, smart watch, etc.
- PC Computer GUI or CLI: Windows, Mac, Linux, etc.
- Laptop Computer GUI or CLI: Windows, Mac, Linux, etc.
- Tablet: Android, Apple, Windows, etc.
- Gaming Console: XBox, Playstation, Steam Box, etc.
- Web Browsers: Chrome, Firefox, etc.
- Virtual Machine: Java, Racket, CLR (Windows .net architecture), etc.

Writing apps for multiple platforms is more possible today than it has been in the past, but on the other hand, there are arguably more viable platforms than in the past, so things aren't that much easier on that front. In this course, we'll develop on one platform: the Linux Command line. More specifically, we'll write code to run on the department's Ubuntu⁸ Linux⁹ server.

For most of you, this is likely to be a new computing environment on two fronts. First, you've probably never seen or maybe even heard of Linux. Even amongst those of you that know about Linux, it's likely you've never really worked at a Command Line Interface (CLI). The command line interface uses no windows, no mouse clicks, just text-based commands. Working at the CLI is roughly equivalent to working at the interactions window of DrRacket. You type commands at a prompt, the computer executes them, and usually prints¹⁰ some results back at the prompt. The CLI pre-dates Graphical User Interfaces¹¹, but is still widely used in many computing environments. It is also usually hiding in the background when GUIs are installed and as such is a viable choice for many projects. In short, the CLI is still alive and kicking and you'll gain a lot from knowing how to work with it. The first thing we'll do in this class is be sure you're able to work and survive at the CLI¹².

⁸ <http://www.ubuntu.com>

⁹ <http://en.wikipedia.org/wiki/Linux>

¹⁰ DrRacket always prints. The CLI does not.

¹¹ GUI

¹² survive = basic working knowledge not ninja CLI hacker skills.

The Criteria of Quality

So we have a problem to solve and a platform on which we'll deploy our computational solution. Or, we have an end in mind and need to start thinking about a means to achieve that end. However, before we start hacking away at some code we should reflect on what it takes to write good code¹³. A craftsman tries to produce quality work and knows how to judge their work for quality. Programs are written in order to be¹⁴:

1. Correct
2. Simple
3. Efficient

Thus the quality of a program¹⁵ can be evaluated on these criteria¹⁶.

If a program does not correctly carry out its intended task, solve its specified problem, then it's not very good. Until the program works correctly, then we don't even need to consider any other measure of quality. This sounds simple enough, but program correctness is a very tricky thing. First off, correct might be subjective and it can be difficult to clearly identify what perfect correctness will be. Furthermore, when we do know what constitutes correctness, it is nearly impossible to guarantee absolute correctness. As such, good programs more often than not exist in a state of *mostly correct*. This is clear from the fact that even the best software needs to fix bugs and update itself. If absolute correctness is unachievable, what then is a programmer to do? First and foremost, programmers need to identify the level of correctness a user can expect and try to guarantee at least that much. This means clearly identifying and documenting the specifications for the program in such a way that a correctness benchmark can be set and *tested*. We must also plan for the unexpected bug to occur and develop code that is easily maintained over the life-time of the program.

Programs should be *simply and elegantly written* so that when bugs appear, you can easily return to the code to fix them¹⁷. Program code is more often read by human beings than it is computers, and as a written document should be judged on those standards. Your code should have structure and style and be easily read by other programmers. As you can imagine, simplicity is subjective and can be hard to evaluate. None the less, there are well accepted styles of programming out there upon which we can choose to evaluate our programs¹⁸. Simplicity is not only a boon to correctness, but a boon to business. Well designed and simply written code is often easy to extend. New features are easier to add to programs that exhibit

¹³ If your goal is to write bad code, you're in the wrong place

¹⁴ In order of importance!

¹⁵ and its programmer

¹⁶ We could be more detailed, but odds are if you meet these criteria, any other more detailed criteria will be covered as well. Alternatively, you should learn to work around these simple goals and only after you've gain experience in this realm, explore more involved criteria of quality programs.

¹⁷ Keep in mind here, we're not really talking about the small programs you've written so far. Yes, even those projects from last semester are small by program standards. We're talking about millions of lines of code.

¹⁸ choose a normal when no universally accepted normal can be found

well established metrics of good design¹⁹. Simple code is also often faster as it cuts out unnecessary program logic and avoid repetition, and when its not fast enough, its simplicity makes it easier to reason about and thereby easier to optimize. And this brings us to goal number three, efficiency.

Sometimes when you focus on writing correct and simple code, you get a program that performs as well as it needs to on the target system. On the other hand, your program may often run too slowly or use up too much memory. Put another way, it may make inefficient use of the computation resources provided to it by the computer. When this is the case²⁰ we must optimize the resource usage of our program. This often means trying to make it run fast, i.e. make better use of the CPU cycles, or use less memory. Of the three programming goals we'll be looking at, this is the newest one for you to think about. We'll need to learn how programmers talk about efficiency and how we measure the efficiency of our code.

Putting this all together. We have a problem, an idea of how to model and solve that problem computationally, and criteria for evaluating the computational solution. Now we need to choose the right tools for the job²¹.

The Tools

DrRacket is a one-stop shop for all your programming tools needs. Such programs are called *Integrated Development Environments*²². You will not be using an IDE in this course²³. Instead we'll look at an established set of industry tools and learn the basics of making them work together. This deconstructed view of the programming tool chain will hopefully give you a better appreciation for the tools that are out there and the different systems tucked away inside IDEs like DrRacket. All the tools you'll be using see wide use today and are viable options of program development.

The Essentials

At a minimum you need a **programming language**²⁴, **text editor to write the code** and either an **interpreter** to execute the code or a **compiler** to build an executable program for your platform. You know what a programming language is and that we'll be using C++ in this class, so we'll focus on the other two tools.

Text editors do what the name implies, they edit text. They do not process words. The difference is that text editors don't really get into presentation details and most importantly do not encode the text in anything other than a plain text encoding²⁵ Windows Notepad is a

¹⁹ the kinds of things the HtDP design recipe gives you

²⁰ and only when this is the case!

²¹ Platform often dictates or restricts your choice of tools. At which point, you have a new problem... building the tool or platform you want!

²² IDE

²³ we'll come back to IDEs in COMP220 and COMP210

²⁴ yes. languages are tools

²⁵ Probably ASCII. Possibly Unicode.

text editor, but not particularly well suited for programming. Good text editors for programming are programmed with information about the language you're using and provide help and cues to ease the task of writing code. DrRacket's definitions window color coded text, helped match parenthesis, fixed indentation to meet Racket style, and much more. These are the types of things we need our text editor to do. For this class, you'll learn to work with GNU²⁶ Emacs²⁷. Other options you might explore include²⁸:

- Vim²⁹
- Sublime Text³⁰
- Notepad++³¹
- Atom ³²

Feel free to explore these and other options, but Emacs is the only supported text editor for this course³³.

Interpreters are installed on the platform and can read and execute code on a line-by-line basis. They run the program as they read the code. On the other hand, a compiler translates the code to another format, typically a fully executable file³⁴. These days it's not uncommon to see a combination of the two. A just-in-time³⁵ compiler will interpret some code but compile performance critical code for faster execution. Racket uses a JIT system. In this class we'll use a traditional compiler system, namely the GNU GCC³⁶ compiler g++. Other notable C++ compilers are:

- LLVM and clang ³⁷
- Visual Studio and CLR ³⁸

Programs quickly grow to involve multiple files. The CLI compiler we'll be using is pretty good at helping you compile all those files quickly with a minimum number of commands. However, it is more common to use a special program building tool to manage the complexity of the compilation process. The build tool *make*³⁹ is a very flexible tool and widely used to manage the problem of making programs. Make is itself a special purpose programming language and interpreter. To use make we write a small file called *Makefile* that the command make then reads and interprets. So, unless you like to enter ten commands when one will do, make is really awesome.

Tools for Correctness

A language, an editor, and a compiler⁴⁰ will get you to a working program. Now we need to address the correctness of that program.

²⁶ <http://www.gnu.org/>

²⁷ <http://www.gnu.org/software/emacs/>

²⁸ some of these are platform dependent

²⁹ <http://www.vim.org/>

³⁰ <http://www.sublimetext.com/>

³¹ <http://notepad-plus-plus.org/>

³² <https://atom.io/>

³³ Don't expect answers to non-Emacs questions

³⁴ or something that's interpreted

³⁵ JIT

³⁶ <http://gcc.gnu.org/>

³⁷ <http://clang.llvm.org/>

³⁸ <http://msdn.microsoft.com/en-us/vstudio/hh386302>

³⁹ <http://www.gnu.org/software/make/manual/make.html>

⁴⁰ and build management system

Good languages often provide you with language features specifically designed to help you write correct code⁴¹. But, practiced programmers also make use of several tools for helping reach their correctness goals. The most common are:

- compilers
- debuggers
- memory checkers
- code testing frameworks

Compilers are the first line of defense. Basic compilers will catch deviations from the language grammar⁴². Unlike your professors, the computer does not⁴³ infer your intentions from your code. In addition to guaranteeing the grammatical correctness of your code, a good compiler will also warn you when you do something that might lead to problems. We'll also see that compilers can effectively annotate our code such that other tools can more effectively analyze it. In particular, g++ can add special flags to the finished product that enable debuggers and profilers to give us better reports about our program's execution behavior.

Grammatical correctness is a pretty weak level of correctness. Every programmer has written a program that compiles and runs but produces incorrect results⁴⁴ or crashes at run-time⁴⁵. A *debugger* allows programmers to step through program execution one step⁴⁶ at a time while keeping an eye on program data. DrRacket had a stepper that allowed for this. For stepping through and debugging our C++ programs, we'll explore the GNU debugger **gdb**⁴⁷.

Run-time errors can often be the result of running afoul of the allowed usage of the computer's memory system. To correct these mistakes, we use programs that observe the memory usage patterns of our program and generate detailed reports of where something goes awry. The standard tool for this in Linux, the tool we're going to use, is *memcheck*⁴⁸. The memcheck tool is a part of the *Valgrind*⁴⁹ family of code analysis instruments.

Where all the previously discussed tools were programs in their own right, testing frameworks are just libraries of code written to more easily enable standard program testing regimes. In COMP160 you learned to do *unit-tests* and we'll continue to use them in this course. These tests look at individual units of the program and test for expected functionality and behavior. In this class we'll make use of a C++ unit testing framework call *gtest*⁵⁰ developed by Google.

⁴¹ assertions and exceptions are two examples

⁴² syntax errors

⁴³ and should not

⁴⁴ logic errors

⁴⁵ run-time errors

⁴⁶ or programmer specified skips

⁴⁷ <https://www.gnu.org/software/gdb/>

⁴⁸ <http://valgrind.org/docs/manual/mc-manual.html>

⁴⁹ <http://valgrind.org/>

⁵⁰ <https://code.google.com/p/googletest/wiki/Documentation>

Tools for Simplicity

There's a general lack of tangible tools for simplicity⁵¹ checking really. The best tool to check for sufficient simplicity is your fellow programmer. Different programming communities often agree upon what good, simply written coding style looks like. These stylistic guidelines ensure that code looks and reads consistently within the community and is thereby simple to the members of the community. So, one of the best things you can do is have your code peer-reviewed⁵² for its style⁵³. We'll adopt some basic style guides for this class and will go over them as we learn C++. In the meantime, you should take a look at what professional style guidelines are like. Google has their C++ style guide published on the web along with style guides for other languages they use⁵⁴. There are programs called *linters* that scan for suspicious looking code, but we won't play with them in this class. Google has an in-development linter/style checking tool called *cpplint*. You might check it out, it's with their style guides.

⁵¹ in the sense that we're using this word

⁵² We need a writing center for code!

⁵³ Don't have your peers write and debug your code for you. That's called academic dishonesty!

⁵⁴ <https://code.google.com/p/google-styleguide/>

Tools for Efficiency

Once the code works and looks good, it's often time to try and speed it up or lower its memory footprint. In most cases, there are no tools that can make your code more efficient for you. You'll need to do your own optimization. Smart programmers involve real data about program performance in their optimization process⁵⁵, so many of our tools are used to gather data that allows us to make informed decisions about optimizing our code.

⁵⁵ Check out: <https://www.facebook.com/notes/facebook-engineering/the-mature-optimization-handbook/10151784131623920>

- Compilers
- Mathematical analysis
- Memory system profilers
- CPU profilers

There is one tool that will auto-magically make⁵⁶ your code faster: the compiler. Modern compilers can carry out basic to sophisticated transformations on common code patterns in order to improve code performance. This is wonderful as we generally just need to focus on big picture logic and not low-level optimization details. However, this process effectively re-writes your code, making it harder to debug so compiler optimizations are often something we don't introduce until we're confident that our program is correct enough. Compiler's can only do so much for you though. If your code is inherently inefficient, then it's not going to fix that for you. This means we

⁵⁶ or attempt to make

need to make good, efficient coding choices before we even turn the compiler loose.

Using a standard form of mathematical analysis⁵⁷, we can guarantee⁵⁸ the worst case behavior of our code under some fairly reasonable assumptions. Once we know we've made sound algorithmic decisions and have acceptable upper-bounds on program efficiency, then we must delve down into reality as our assumptions are reasonable but simplify some key details. To see what happens to our code on hardware we must run it and use a *profiler* to gather performance metrics about its execution. For the types of programs we're looking at, we need to know how efficiently our program makes use of the CPU and of memory system.

Valgrind provides us with a CPU profiler called *callgrind*⁵⁹. This profiler attempts to count how often each procedure is called and where in the code it's called. From this we can begin to understand what code is running most often and where we can get the most bang for our optimization buck⁶⁰. It should be noted that to profile the CPU we often count the things that are executed and not how long they take to execute. We'll come back to this. For now, you should think about why that might be a good idea.

Valgrind also provides us with a memory system profiler called *cachegrind*⁶¹ and another called *massif*⁶². These tools let us look at different parts of the memory system⁶³ and determine how often we're using them and if we're using them efficiently.

Other Tools

The final, commonly used tool, that we're likely to play with is a Version Control System⁶⁴ called *git*⁶⁵. You hopefully have picked up on the possibility that real programs have long life spans. You write some code then fix somethings and optimize others. In professional settings in particular, it is important to keep track of code that results in the last, stable piece of software you developed. Version control systems effectively let you take snapshots of your code and then do things like jump back to a previous snapshot or merge new code with an existing snapshot. They also enable easy off-sight backup in case the computer you're working on goes kaput and you lose your code⁶⁶.

Tool Wrap-up

That's a lot of tools. We'll just barely scratch the surface of what most of them can do. Our goal is not to master these tools but to recognize that they're there and how they are used to develop better programs. If you come at this from the other direction, their very existence and

⁵⁷ Asymptotic analysis or "Big-Oh"

⁵⁸ as in mathematically prove!

⁵⁹ <http://valgrind.org/docs/manual/cl-manual.html>

⁶⁰ Make the common case fast

⁶¹ <http://valgrind.org/docs/manual/cg-manual.html>

⁶² <http://valgrind.org/docs/manual/ms-manual.html>

⁶³ the cache and heap respectively

⁶⁴ VCS

⁶⁵ <http://git-scm.com/>

⁶⁶ checkout <http://github.com>

functionality sheds light on what matters to practiced programs. Programmers encounter problems with developing good⁶⁷ code and these are the programs they developed to solve their problems. So even if the types of programs we write in this and other classes don't really need these tools, we can rest assured that some day we'll bump in to the exact kinds of problems these tools were developed to solve. With that in mind, let's revisit the tools we'll be putting to use in this class:

⁶⁷ correct, simple, and efficient

1. Platform: Linux CLI
2. Language: C++
3. Text Editor: Emacs
4. Compiler: g++
5. Build Automation: make
6. Unit Testing Framework: gtest
7. Debugger: gdb
8. Memory Checker: valgrind memcheck
9. Efficiency Analysis: mathematics. asymptotic analysis.
10. CPU Profiler: valgrind callgrind
11. Memory Cache Profiler: valgrind cachegrind
12. Stack and Heap Profiler: valgrind massif
13. VCS: git

Finally, one of our best tools in our tool box is our fellow programmer. Programming is, more often than not, something done by a community of like minded individuals setting out to solve some problems. Feedback from your programming community can really help improve your code and make you a better programmer.